

# STAT 745: Homework 2

Ian Effendi (iae2784@rit.edu)

9/16/2021

## Contents

Overview . . . . .	1
Preprocess . . . . .	2
Contents of <code>Liver.txt</code> . . . . .	2
Clean <code>Liver.txt</code> . . . . .	2
Classification . . . . .	5
Logistic Regression . . . . .	5
Model Analysis . . . . .	5
Misclassification Error Analysis . . . . .	6
Crossvalidation . . . . .	6
$k = 10$ -fold Crossvalidation . . . . .	6
$k = 3$ -fold Crossvalidation . . . . .	6
Session Information . . . . .	6

```
# Required packages
require(knitr)
require(rmarkdown)
require(tidyverse) # readr, tibble, ggplot2
```

## Overview

This section provides an overview of the document.

For this assignment, we will be exploring the `Livers.txt` data set, detailing various liver and blood conditions of 345 men, with the goal of predicting liver disorder severity.

## Preprocess

In this section, I clean up the `Liver.txt` data source and output an intermediary data file to make it easier to perform classification with the data.

### Contents of `Liver.txt`

1. Read in the data from the text file "`Liver.txt`".

The data is stored in a raw text file with comma-separated values. We can use `readr::read_csv()` on this file as-is to get our data into R. Additionally, the `read_csv()` accepts a `col_types` parameter that allows us to designate the types for each of the features.

```
# Read raw text file into a tibble.
Liver.txt <- as_tibble(read_csv(
  # File is located in the data/raw folder.
  file = "data/raw/Liver.txt",
  # First row is not a header row, so we should set `col_names` to FALSE.
  col_names = FALSE,
  # The column type choices are described below.
  col_types = cols(
    # 1 through 5 are quantitative.
    X1 = col_integer(),
    X2 = col_integer(),
    X3 = col_integer(),
    X4 = col_integer(),
    X5 = col_integer(),
    # No. of alcoholic beverages drunk.
    X6 = col_double(),
    # Severity group of the liver condition.
    X7 = col_factor(),
  )
))

# Preview the first 10 rows from the data set.
Liver.txt
#> # A tibble: 345 x 7
#>       X1     X2     X3     X4     X5     X6 X7
#>   <int> <int> <int> <int> <int> <dbl> <fct>
#> 1    85    92    45    27    31    0    1
#> 2    85    64    59    32    23    0    2
#> 3    86    54    33    16    54    0    2
#> 4    91    78    34    24    36    0    2
#> 5    87    70    12    28    10    0    2
#> 6    98    55    13    17    17    0    2
#> 7    88    62    20    17     9    0.5  1
#> 8    88    67    21    11    11    0.5  1
#> 9    92    54    22    20     7    0.5  1
#> 10   90    60    25    19     5    0.5  1
#> # ... with 335 more rows
```

We can see the tibble contains 345 rows and a total of 7 columns, which matches our expectations from the assignment details.

### Clean `Liver.txt`

The assignment document provides some information about the features in the data set. We know that:

- Variables 1 through 5 are quantitative results of blood tests sensitive to liver disorders.
  - Inference: These are stored as whole integers in the original file, so a `<int>` type seems appropriate.
  - Concern: We don't really know what units these are stored with; if they are percentages, some level of scaling or data transformation may have been appropriate.
- Variable 6 is the number of alcoholic beverages drunk per day.
  - Inference: This is stored as a fractional numeric, so despite being described as something that I would assume is a discrete integer, we'll store this as a `<dbl>` as well.
- Variable 7 is the response, demarcating the liver malfunctioning group.
  - Inference: This is a factor. Given this is a classification problem, we should store this as a categorical `<fctr>` type.

```
# Normalize column names to lowercase.
Liver.1 <- rename_with(Liver.txt, tolower)

# Rename some column names to be more descriptive.
# Uses `dplyr::rename(.data, new.name = old.name)`

# x1 through x5 will have similar naming convention
# as they are blood test results sensitive to liver conditions.
Liver.1 <- rename(Liver.1, blood.1 = x1)
Liver.1 <- rename(Liver.1, blood.2 = x2)
Liver.1 <- rename(Liver.1, blood.3 = x3)
Liver.1 <- rename(Liver.1, blood.4 = x4)
Liver.1 <- rename(Liver.1, blood.5 = x5)

# x6 -> drinks
Liver.1 <- rename(Liver.1, drinks = x6)

# x7 -> severity
Liver.1 <- rename(Liver.1, severity = x7)

# Preview changes to column names.
names(Liver.1)
#> [1] "blood.1" "blood.2" "blood.3" "blood.4" "blood.5"
#> [6] "drinks"   "severity"
```

Now that the fields are named appropriately, I want to finish cleaning the data up by encoding the `severity` factor as a binary field.

Table 1: Severity Encoding Map

Status	Value	Code
severe	1	1
not severe	2	0

```
# Encode severity == 1 as 1, All other values as 0.
Liver.2 <- Liver.1 %>%
  mutate(severity = as.integer(ifelse(severity == 1, 1, 0)))

# Preview the Liver tibble.
Liver.2
#> # A tibble: 345 x 7
#>   blood.1 blood.2 blood.3 blood.4 blood.5 drinks severity
#>   <int>   <int>   <int>   <int>   <int>   <dbl>   <int>
```

```

#>   1      85      92      45      27      31      0      1
#>   2      85      64      59      32      23      0      0
#>   3      86      54      33      16      54      0      0
#>   4      91      78      34      24      36      0      0
#>   5      87      70      12      28      10      0      0
#>   6      98      55      13      17      17      0      0
#>   7      88      62      20      17      9      0.5      1
#>   8      88      67      21      11      11      0.5      1
#>   9      92      54      22      20      7      0.5      1
#>  10      90      60      25      19      5      0.5      1
#> # ... with 335 more rows

# Write the data to the `data/processed/` directory as `Liver.Rds`
write_rds(Liver.2, "data/processed/Liver.Rds")

```

Instead of using the `save()` function (which would save our entire environment), we'll save the single modified `Liver.2` object with `readr::write_rds()` wrapper for `saveRDS()`.

This method will allow us to import our intermediate data object into a new variable rather than dumping loose data objects into the environment, when using the `readRDS()` wrapper, `readr::read_rds()`.

## Classification

2. Predict Variable 7 using logistic regression (using  $p = P(\text{Group } 1)$ ) with the remaining variables as predictors.

```
# Read the processed intermediate file.
Liver <- read_rds("data/processed/Liver.Rds")
Liver[1,]
#> # A tibble: 1 x 7
#>   blood.1 blood.2 blood.3 blood.4 blood.5 drinks severity
#>   <int>   <int>   <int>   <int>   <int>   <dbl>   <int>
#> 1      85      92      45      27      31      0      1
```

## Logistic Regression

*severity* can only take one of two values: 0 or 1. We denote  $p = P(\text{severity} = 1)$  and we will fit a logistic regression model:

$$\ln\left[\frac{p}{1-p}\right] = \beta_0 + \beta_1(\text{blood.1}) + \cdots + \beta_5(\text{blood.5}) + \beta_6(\text{drinks})$$

We will predict *severity* = 1 if  $p > p_0$  and as 0 otherwise.

```
# Define and fit a logistic regression model.
fit.1 <- glm(severity ~ ., data=Liver, family="binomial")
```

## Model Analysis

```
# Summary fit function from class.
Summary.fit <- function(fit.obj) {
  mat <- summary(fit.obj)$coef
  mat[,3] <- round(mat[,3], 2)
  mat[,4] <- round(mat[,4], 2)
  return(mat)
}

# Show results.
Summary.fit(fit.1)

fit.1$residuals <- as_tibble(residuals.glm(fit.1))
fit.1$residuals
#>           Estimate Std. Error z value Pr(>|z|)
#> (Intercept) -5.99025807 2.685457986  -2.23    0.03
#> blood.1      0.06398364 0.029644744   2.16    0.03
#> blood.2      0.01952514 0.006759465   2.89    0.00
#> blood.3      0.06410602 0.012299568   5.21    0.00
#> blood.4     -0.12319809 0.024273023  -5.08    0.00
#> blood.5     -0.01894716 0.005601748  -3.38    0.00
#> drinks      0.06808016 0.040379564   1.69    0.09
#> # A tibble: 345 x 1
#>   value
#>   <dbl>
#> 1  1.09
#> 2 -1.22
#> 3 -1.05
#> 4 -1.13
```

```
#> 5 -0.521
#> 6 -1.08
#> 7 1.22
#> 8 0.874
#> 9 1.26
#> 10 1.12
#> # ... with 335 more rows
```

## Misclassification Error Analysis

3. Minimize the total error of misclassification by changing the threshold for  $p$ . Show the results graphically, as we did in class, and also show the optimal values, including the misclassification table for that optimal case.

```
# Make the prediction needed with threshold set to  $p = 0.5$ .
p.threshold <- 0.5
fit.1.prediction <- (fit.1$fitted > p.threshold)

# Get the misclassification table.
(Classif.table <- table(fit.1.prediction, Liver$severity))

# Calculate the error rate.
Error.rate.fit <- function(mat) {
  # Assume a 2 by 2 table.
  (mat[1,2] + mat[2,1]) / sum(mat)
}

# Misclassification error rate:
Error.rate.fit(Classif.table)

# Normal 'at random' error rate:
sum(Liver$severity==0)/length(Liver$severity)
#>
#> fit.1.prediction  0  1
#>          FALSE 165  67
#>          TRUE   35  78
#> [1] 0.2956522
#> [1] 0.5797101
```

## Crossvalidation

$k = 10$ -fold Crossvalidation

4. Repeat Step 3 with 10-fold CV. Use 100 rounds. Calculate means and their standard errors of the minimized total error of misclassification.

$k = 3$ -fold Crossvalidation

5. Repeat Step 3 with 3-fold CV. Use 100 rounds. Calculate means and their standard errors of the minimized total error of misclassification.

## Session Information

This report was generated by the `vignettes/hw2_notebook.Rmd` file. The setup chunk for this document sets the root directory to the project root directory using the `rprojroot` package; all file paths are relative to the project root.

```

#> R version 4.1.1 (2021-08-10)
#> Platform: x86_64-w64-mingw32/x64 (64-bit)
#> Running under: Windows 10 x64 (build 19043)
#>
#> Matrix products: default
#>
#> locale:
#> [1] LC_COLLATE=English_United States.1252
#> [2] LC_CTYPE=English_United States.1252
#> [3] LC_MONETARY=English_United States.1252
#> [4] LC_NUMERIC=C
#> [5] LC_TIME=English_United States.1252
#>
#> attached base packages:
#> [1] stats      graphics  grDevices  utils      datasets
#> [6] methods    base
#>
#> other attached packages:
#> [1] forcats_0.5.1  stringr_1.4.0  dplyr_1.0.7
#> [4] purrr_0.3.4    readr_2.0.1    tidyr_1.1.3
#> [7] tibble_3.1.4   ggplot2_3.3.5  tidyverse_1.3.1
#> [10] rmarkdown_2.10 rprojroot_2.0.2 knitr_1.34
#>
#> loaded via a namespace (and not attached):
#> [1] tidyselect_1.1.1 xfun_0.25      haven_2.4.3
#> [4] colorspace_2.0-2 vctrs_0.3.8    generics_0.1.0
#> [7] htmltools_0.5.2  yaml_2.2.1     utf8_1.2.2
#> [10] rlang_0.4.11     pillar_1.6.2   glue_1.4.2
#> [13] withr_2.4.2      DBI_1.1.1      bit64_4.0.5
#> [16] dbplyr_2.1.1     modelr_0.1.8   readxl_1.3.1
#> [19] lifecycle_1.0.0  munsell_0.5.0  gtable_0.3.0
#> [22] cellranger_1.1.0 rvest_1.0.1    evaluate_0.14
#> [25] tzdb_0.1.2       fastmap_1.1.0  parallel_4.1.1
#> [28] fansi_0.5.0      highr_0.9      broom_0.7.9
#> [31] Rcpp_1.0.7       scales_1.1.1   backports_1.2.1
#> [34] vroom_1.5.4      jsonlite_1.7.2 bit_4.0.4
#> [37] fs_1.5.0         hms_1.1.0      digest_0.6.27
#> [40] stringi_1.7.4    grid_4.1.1     cli_3.0.1
#> [43] tools_4.1.1      magrittr_2.0.1 crayon_1.4.1
#> [46] pkgconfig_2.0.3  ellipsis_0.3.2 xml2_1.3.2
#> [49] reprex_2.0.1     lubridate_1.7.10 rstudioapi_0.13
#> [52] assertthat_0.2.1 httr_1.4.2     R6_2.5.1
#> [55] compiler_4.1.1

```